

ECE36800 Data structures

Mergesort

Chapter 8 (pp. 335-361)

Lu Su

School of Electrical and Computer Engineering
Purdue University

Overview

- [illegible]

Recursive mergesort

- Mergesort algorithm

- 1) Split array into two subarrays of about the same size
- 2) Sort each subarray recursively (if size > 1)
- 3) Merge the sorted subarrays

- Complexity

- Recursion depth is at most $\log n + 1 = O(\log n)$
- If merging two sorted subarrays takes $O(n)$, mergesort is $O(n \log n)$

Mergesort: algorithm

- Sort n integers $r[0]$ to $r[n-1]$ in ascending order
- Call `mergesort(r, 0, n - 1)`

```
mergesort (r[], lidx, ridx):  
    if lidx ≥ ridx  
        return;  
    // divide  
    mid = (lidx+ridx)/2; // assume no overflow  
    // conquer left subarray  
    mergesort(r, lidx, mid);  
    // conquer right subarray  
    mergesort(r, mid+1, ridx);  
    // merge subarrays  
    merge(r, lidx, mid, ridx)
```

Mergesort

- Another divide-and-conquer algorithm (natural to use recursion)
- Quicksort: “difficult” division, easy combination, preorder

```
pivot_idx = partition(r, lidx, ridx);  
qsort(r, lidx, pivot_idx-1);  
qsort(r, pivot_idx+1, ridx)
```

- Mergesort: easy division, “difficult” combination, postorder







```
mergesort(r, lidx, mid);  
mergesort(r, mid+1, ridx);  
merge(r, lidx, mid, ridx)
```

Mergesort: merge


- Merge sorted subarrays $r[lidx..mid]$ and $r[mid+1..ridx]$ in ascending order
- Assume auxiliary space $tmp[0..n-1]$, where n is the number of elements in the array
- The function `memcpy` copies the contents at second parameter to the location at first parameter; the number of elements copied is specified by third parameter
 - Not exactly the function in `string.h`

```
merge (r[], lidx, mid, ridx):  
    memcpy(&tmp[lidx], &r[lidx], mid-lidx+1);  
    memcpy(&tmp[mid+1], &r[mid+1], ridx-mid);  
    i ← lidx; j ← mid+1;  
    for k ← lidx to ridx  
        if (i > mid) r[k] ← tmp[j++];  
        else if (j > ridx) r[k] ← tmp[i++];  
        else if (tmp[j] < tmp[i])  
            r[k] ← tmp[j++];  
        else r[k] ← tmp[i++];
```

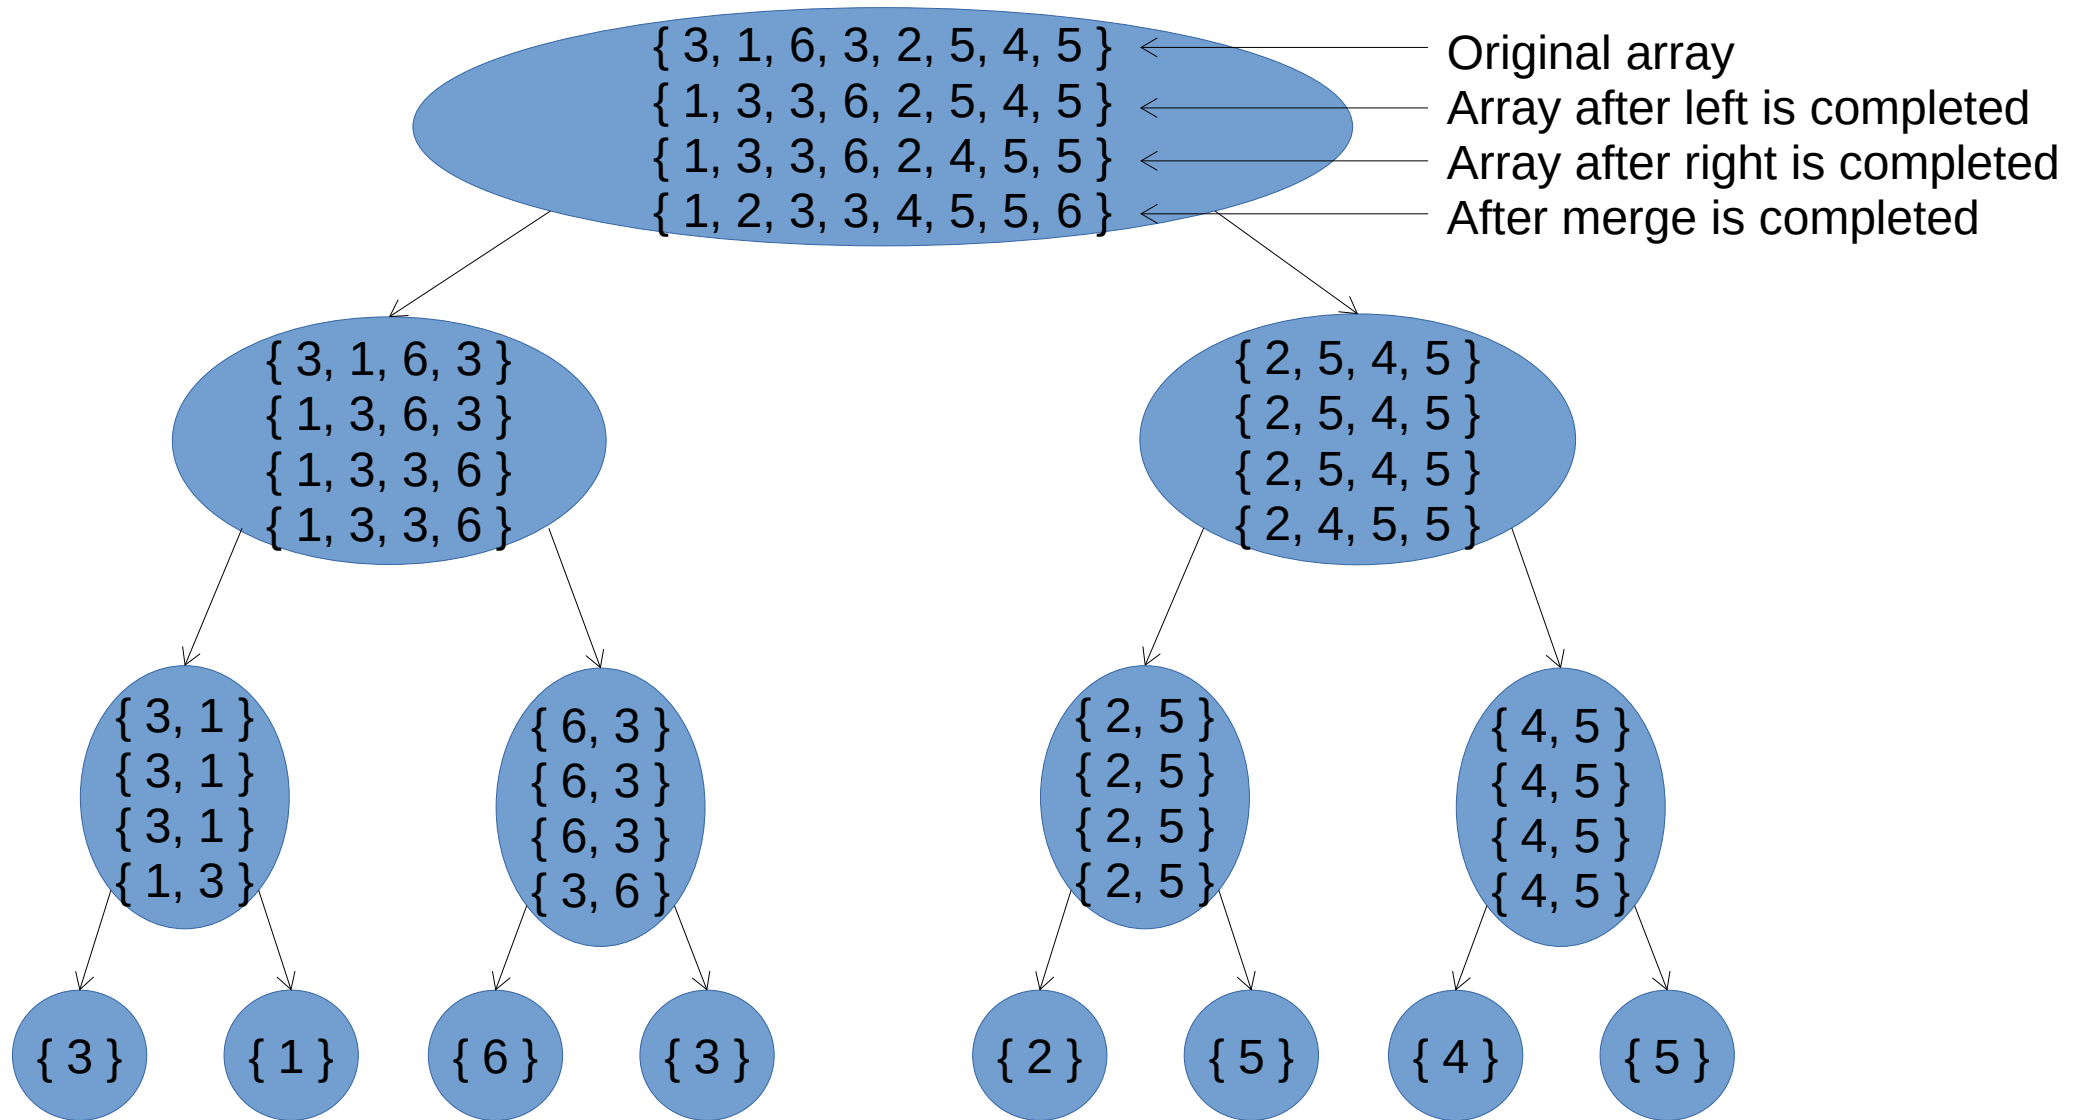
Merge example



Merge example




Mergesort example



Mergesort: Iterative algorithm

- Bottom-up process
- Leaf nodes: Consider original array as n subarrays, each of size 1
- Scan through array performing $n/2$ times, merging of two 1-element arrays to produce $n/2$ sorted subarrays, each of size 2
- Scan through array performing $n/4$ times, merging of two 2-element arrays to produce $n/4$ sorted subarrays, each of size 4
- ...
- Perform merging of two $n/2$ -element arrays to produce the final sorted array of n elements

Iterative mergesort example

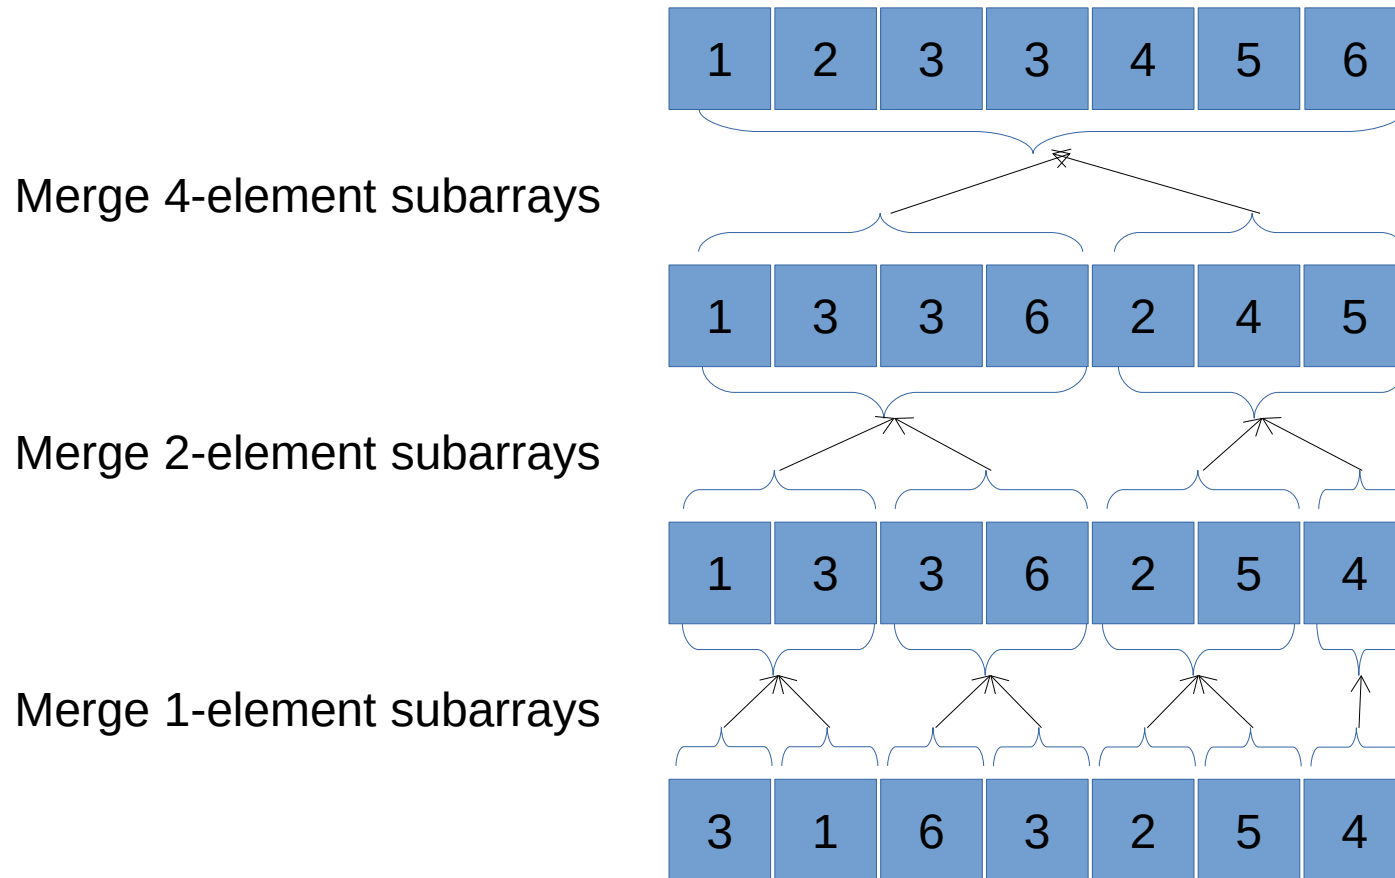


Mergesort: Iterative algorithm

- Iteratively mergesort n integers $r[0..n-1]$ in ascending order

```
iter_mergesort (r[], n):  
    size ← 1  
    while size < n  
        i ← 0  
        while i < n - size  
            merge(r, i, i+size-1,  
                  min(i+2*size-1, n-1));  
            i ← i + 2*size  
        size ← 2*size;
```

Iterative mergesort example



Properties of mergesort

- Memory complexity: $O(n)$, not in-place
 - $O(n)$ copy operations in each pass
 - Main disadvantage to quicksort or heapsort
 - In-place variant exists, but large overhead
- Time complexity: $O(n \log n)$
 - For a sorted array, worse than insertion sort
- Can be stable if merge() is stable
- Easier to implement than quicksort and heapsort
- In practice slower than quicksort, but faster than heapsort
- Advantage: sequential access of data
 - Excellent for external (out-of-core) sorting (e.g. disk)
 - Good cache usage

Bitonic merge

- Inefficiency in merge function: Checks whether indices are out-of-bound
- Use sentinel (as in insertion sort) to avoid checking whether indices are within bounds

- Not stable

```
bsmerge (r[], lidx, mid, ridx):  
    memcpy(&tmp[lidx], &r[lidx], mid-lidx+1);  
    j ← ridx;  
    for k ← mid+1 to ridx  
        tmp[k] ← r[j]; // copy in reverse  
        j--;           // order  
    i ← lidx; j ← ridx;  
    for k ← lidx to ridx  
        if (tmp[j] < tmp[i]) r[k] ← tmp[j--];  
        else r[k] ← tmp[i++];
```

Time-complexity summary

Algorithms	Best	Average	Worst
Insertion	$O(n)$	$O(n^2)$	$O(n^2)$
Bubble	$O(n)$	$O(n^2)$	$O(n^2)$
Quick	$O(n \log n)$	$O(n \log n)$	$O(n^2)$
Simple selection	$O(n^2)$	$O(n^2)$	$O(n^2)$
Heap	$O(n)$	$O(n \log n)$	$O(n \log n)$
Merge	$O(n \log n)$	$O(n \log n)$	$O(n \log n)$

$O(n^2)$ vs. $O(n \log n)$

- How can mergesort, quicksort, and heapsort achieve $O(n \log n)$ when insertion sort and bubble sort run at $O(n^2)$?

- Mergesort, quicksort, and heapsort move elements far distances, correcting multiple inversions (incorrect ordering) at a time
- Insertion and bubble sort correct one inversion at a time

- Why is quicksort worst-case $O(n^2)$ while mergesort has no such problem?

- The choice of pivot determines size of partitions, whereas mergesort cuts array in half every iteration
- Simple selection sort uses the worst pivot in every iteration

- Is $O(n \log n)$ the best we can do?

- Yes, if we use binary comparison on the key values

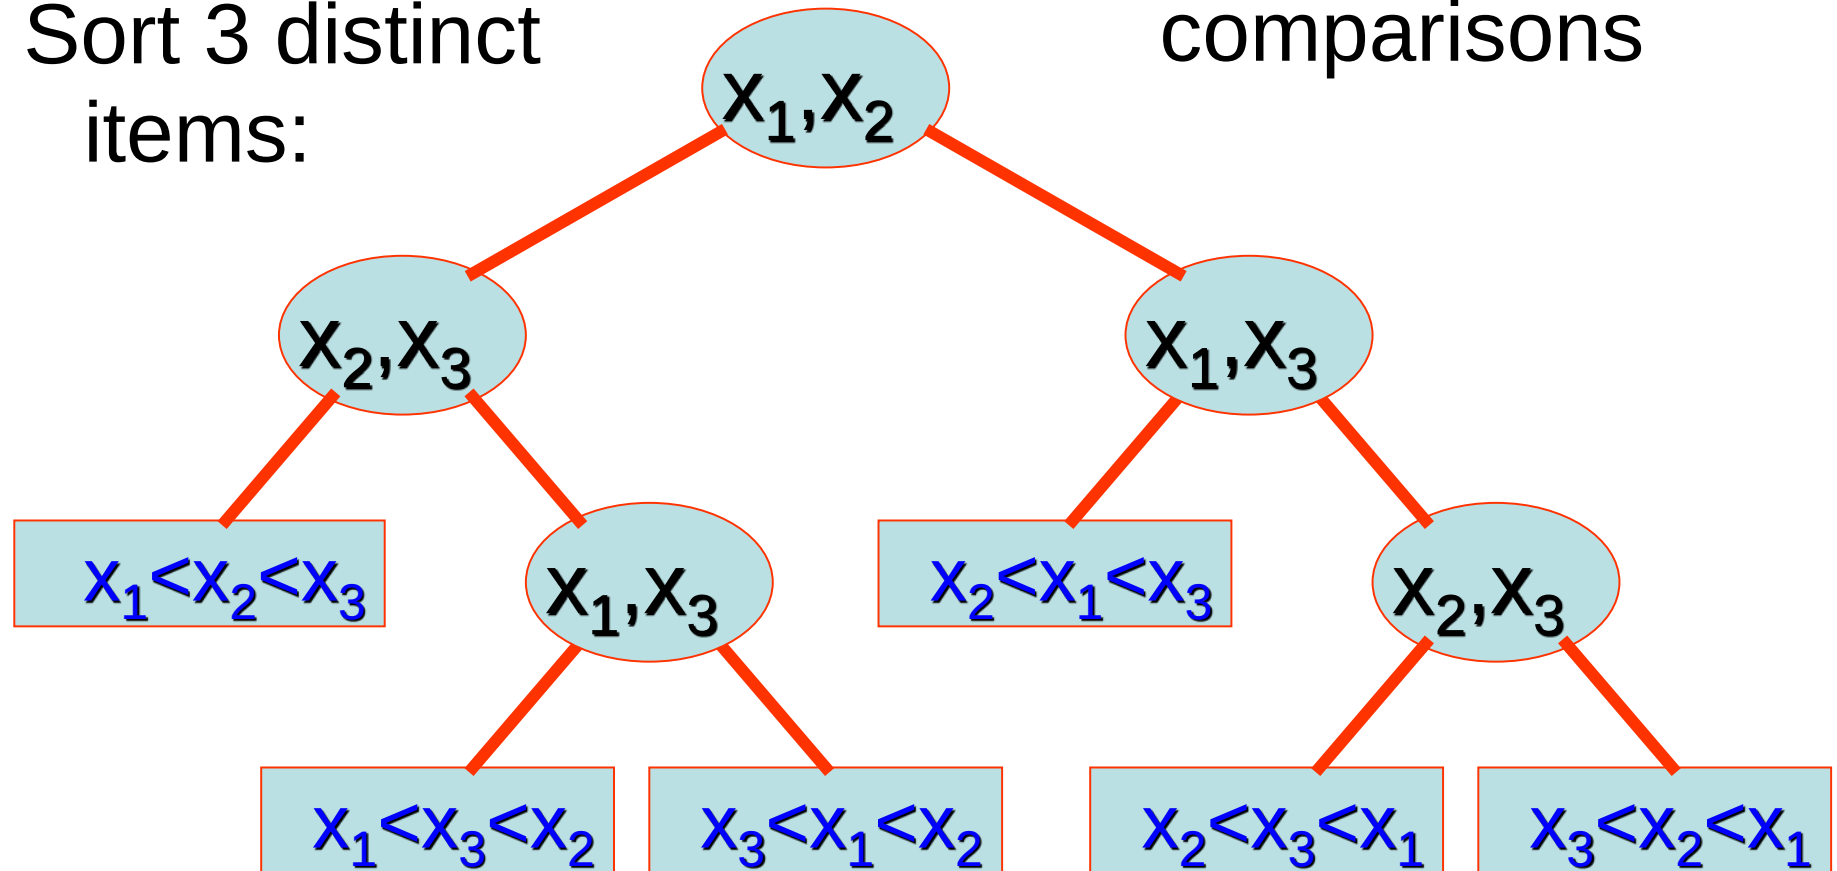
Theoretic lower bound

- n elements can have $n!$ permutations
- A sort decision tree is a binary tree that represents a sorting method based on comparisons
- Each internal node compares two elements
 - Left branch implies first element smaller than second
 - Right branch otherwise
 - No redundant comparisons
- Each leaf nodes corresponds to one of the $n!$ permutations

Binary decision tree for sorting

No redundant
comparisons

Sort 3 distinct
items:



leaf nodes = $n!$

Theoretic lower bound

- A binary decision tree for sorting has $(2^{n!} - 1)$ nodes
- Height of tree is at least $\log n!$
- $n! \geq (n/2)^{(n/2)} \Rightarrow \log n!$ is $\Omega(n \log n)$
- Sorting algorithms that use comparisons must make $\Omega(n \log n)$ comparisons

Quick, Merge, or Heap?

- Quicksort is the fastest sorting algorithm in practice
 - Can perform as badly as $O(n^2)$
 - Requires $O(\log n)$ stack space
- Mergesort is guaranteed to run in $O(n \log n)$
 - Slower than quicksort on average
 - Requires $O(n)$ additional temporary storage
- Heapsort is guaranteed to run in $O(n \log n)$, in-place, and requires no recursions
 - Many real world tests show that heapsort is slower than quicksort (and mergesort) on average
 - Useful for really large problems